


Optimizing PyTrim (Pt. 2)

Using a combination of NumPy and Numba

Content

- Benchmarks
- Implementation
 - Improving performance using pure NumPy
 - Adding Numba
 - Parallelizing
- Conclusion

Overview of simulation times



Improving performance using pure NumPy

- Current version of PyTrim heavily relies on python loops
 - NumPy can't optimize them -> no SIMD optimizations possible
 - Numba can, but has its own limitations (e.g. reduced [Num]Py feature support)

Proposal

- Vectorize everything and rely on more efficient NumPy methods
 - Adding Numba later should be possible and may be beneficial

Improving performance using pure NumPy

- Current state: a maximum of 3 "parallel" arithmetical operations
 - Full simulation for a single ion per iteration
- *Proposal:* as much as possible independent "parallel" arithmetical operations
 - Single simulation step for *all* particles per iteration
 - Parallelization may be possible / beneficial on larger number of particles

```
import numpy as np

a = np.random.rand(5, 3)
b = np.random.rand(5, 3)
# 15 simultaneous operations
a += b
# Time with 100k rows: 0.003s
print(a)
```

DEMO

Improving performance using pure NumPy

Adding Numba

The following methods / signatures did not compile with Numba:

```
np.linalg.norm(axis)
np.ndarray[list, list]
np.ndarray[:, list[list]]
np.put_along_axis()
np.argmin(keepdims)
```

(full list can be found [here](#))

Those parts had to be re-written, sometimes in ugly ways

PyQT integration: Adding Numba prevents the use of Queues to send simulation updates to UI thread


Adding Numba - Parallelization

```
@jit(fastmath=True, parallel=True, nogil=True, cache=True)
def simulate(nion: int, nthreads: int = 16):
    # Initial conditions of the projectile
    # NOTE In reality a check is performed to split data in equal parts
    threads = nthreads
    col_cnt = nion / nthreads
    e = np.full((threads, col_cnt), 50000.0) # energy (eV)
    pos = np.zeros((threads, col_cnt, 3)) # position (A)
    dir = np.zeros((threads, col_cnt, 3))
    dir[:, :, 2] = 1.0 # direction (unit vector)
    is_inside = np.full((threads, col_cnt), True)

    for i in prange(threads):
        trajectory.trajectories(pos[i], dir[i], e[i], is_inside[i])

    pos = pos.reshape((nion, 3))
    dir = dir.reshape((nion, 3))
    e = e.flatten()
    is_inside = is_inside.flatten()
    count_inside = np.count_nonzero(is_inside)
    mean_z = pos[is_inside, 2]
    std_z = np.std(mean_z)
    mean_z = np.mean(mean_z)
```

Overview of simulation times



Conclusion

- Provided NumPy implementation brings a significant performance boost
 - Additional 2x possible by reducing precision from `float64` -> `float32`
- Adding Numba brings a slight improvement, but also significant limitations and less readable code
 - Limited [Num]Py feature set
 - Increased codebase size / complexity
- *UI integration options*
 - Vanilla NumPy implementation (in background) notifying about finished simulations (processing in main / other thread)
 - Numba with data splitting, but with Python threads

Questions?